

# Homework 4

Youhui Ye

10/5/2020

## Problem 1

```
set.seed(1256)
theta <- as.matrix(c(1,2),nrow=2)
X <- cbind(1,rep(1:10,10))
h <- as.vector(X%*%theta+rnorm(100,0,0.2))
## initialization for loop
m <- nrow(X)
theta0old <- Inf
theta1old <- Inf
theta0new <- 0
theta1new <- 0
alpha <- 0.05
tol <- 1e-05
while ((abs(theta0old - theta0new) > tol) || (abs(theta1old - theta1new) > tol)) {
  theta0old <- theta0new
  theta1old <- theta1new
  pred <- X[,1] * theta0old + X[, 2] * theta1old
  theta0new <- theta0old - alpha / m * sum(pred - h)
  theta1new <- theta1old - alpha / m * sum((pred - h) * X[,2])
}
cat("Tolerance: ", tol)

## Tolerance: 1e-05
cat("Step size: ", alpha)

## Step size: 0.05
cat("estimated theta0: ", theta0new)

## estimated theta0: 0.968629
cat("estimated theta1: ", theta1new)

## estimated theta1: 2.001698
m1 <- lm(h ~ 0 + X)
m1$coefficients
```

```
##          X1          X2
## 0.9695707 2.0015630
```

The answer is very close to what “lm” function gives.

user	system	elapsed
20.56	11.47	72362.59

## Problem 2

### Part a

```
## set up
theta0s <- seq(0, 2, length.out = 100)
theta1s <- seq(1, 3, length.out = 100)
thetas <- expand.grid(theta0s, theta1s)
## wrap up the function to implement parallel computing
init <- thetas[1,]
my_gradient_descent <- function(init, X, h) {
  ## set up
  m <- 100
  alpha <- 1e-07
  tol <- 1e-09
  theta0old <- Inf
  theta1old <- Inf
  theta0new <- as.numeric(init[1])
  theta1new <- as.numeric(init[2])
  ## iteration time
  i <- 0
  while ((abs(theta0old - theta0new) > tol) || (abs(theta1old - theta1new) > tol)) {
    theta0old <- theta0new
    theta1old <- theta1new
    pred <- X[,1] * theta0old + X[, 2] * theta1old
    theta0new <- theta0old - alpha / m * sum(pred - h)
    theta1new <- theta1old - alpha / m * sum((pred - h) * X[,2])
    i <- i + 1
    if(i > 5e06) break
  }
  print("yes")
  return(c(theta0new, theta1new, i))
}
## Using parallel computing
library(parallel)
cores <- detectCores() - 1
## log.txt to trace how many runs already completed
cl <- makeCluster(cores, outfile = "log.txt")
clusterExport(cl, "X")
clusterExport(cl, "h")
system.time(result <- parApply(cl, thetas, 1, my_gradient_descent, X, h))
stopCluster(cl)

min(result[3,]) - 1

## [1] 633612
max(result[3,]) - 1

## [1] 5e+06
```

```
## beta_0 estimation
mean(result[1,])
```

```
## [1] 0.9966097
```

```
sd(result[1,])
```

```
## [1] 0.5197305
```

```
## beta_1 estimation
mean(result[2,])
```

```
## [1] 1.997679
```

```
sd(result[2,])
```

```
## [1] 0.07465446
```

## Part b

I do not think it is a good stopping rule. The problem is that the loop may never stop when it reaches a local

## Part c

The algorithm is better used for smooth function. When the function value varies vigorously and the step size is too small, it is very likely to get stuck in a local minimum. Also, there is a lot of work in choosing start values.

## Problem 3

```
beta_hat <- solve(t(X) %*% X) %*% t(X) %*% h
```

Least square estimation is subject to find minimizer of  $(y - X\hat{\beta})'(y - X\hat{\beta})$ . If we take derivative with respect to this function, we have  $-2X'(y - X\hat{\beta}) = 0$ . The answer is exactly  $(X'X)^{-1}X'y$ .

## Problem 4

```
set.seed(12456)
G <- matrix(sample(c(0,0.5,1),size=16000,replace=T),ncol=10)
R <- cor(G) # R: 10 * 10 correlation matrix of G
C <- kronecker(R, diag(1600)) # C is a 16000 * 16000 block diagonal matrix
id <- sample(1:16000,size=932,replace=F)
q <- sample(c(0,0.5,1),size=15068,replace=T) # vector of length 15068
A <- C[id, -id] # matrix of dimension 932 * 15068
B <- C[-id, -id] # matrix of dimension 15068 * 15068
p <- runif(932,0,1)
r <- runif(15068,0,1)
C <- NULL #save some memory space
```

## Part a

```
object.size(A)
```

```
## 112347224 bytes
```

```
object.size(B)
```

```
## 1816357208 bytes
```

```
## around 15 mins
```

```
system.time(y <- p + A %*% solve(B) %*% (q - r))
```

```
## user system elapsed
```

```
## 748.65 12.80 773.78
```

## Part b

I think the inverse of matrix can be completed independently and in a different way. And the multiplication of A and inverse of B can be completed faster.

## Part c

```
## Use c++ code to speed up computing speed
```

```
require(RcppEigen)
```

```
require(inline)
```

```
## matrix multiplication by c++
```

```
txt <- "  
using Eigen::Map;  
using Eigen::MatrixXd;  
using Rcpp::as;  
NumericMatrix tm22(tm2);  
NumericMatrix tmm(tm);  
const MatrixXd ttm(as<MatrixXd>(tmm));  
const MatrixXd ttm2(as<MatrixXd>(tm22));
```

```
MatrixXd prod = ttm*ttm2;  
return(wrap(prod));  
"
```

```
mul_cpp <- cxxfunction(signature(tm="NumericMatrix",  
                                tm2="NumericMatrix"),  
                      plugin="RcppEigen",  
                      body=txt)
```

```
## matrix inversion by c++
```

```
txt2 <- "  
using namespace Rcpp;  
using Eigen::Map;  
using Eigen::VectorXd;  
using Eigen::MatrixXd;  
typedef Map<MatrixXd> MapMatd;  
const MapMatd tmm(as<MapMatd>(tm));  
const MatrixXd tmm_inv = tmm.inverse();  
return( wrap(tmm_inv));"
```

```
solve_cpp <- cxxfunction(signature(tm="NumericMatrix"),  
                        plugin="RcppEigen",  
                        body=txt2)
```

```
## around 13.25 mins, slightly faster
system.time({
  B_inv <- solve_cpp(B)
  y <- p + mul_cpp( mul_cpp(A, B_inv), as.matrix(q - r))
})
```

Using R built-in function takes up 15 minutes to calculate the  $\theta_{\text{hat}}$ . However, if I use RcppEigen library, it takes 13 minutes, which is faster.

## Problem 5

a

```
compute_proportion <- function(vec) {
  sum(vec) / length(vec)
}
```

b

```
set.seed(12345)
P4b_data <- matrix(rbinom(10, 1, prob = (31:40)/100), nrow = 10, ncol = 10, byrow = FALSE)
```

c

```
## calculate by row
apply(P4b_data, 1, compute_proportion)
```

```
## [1] 1 1 1 1 0 0 0 0 1 1
```

```
## calculate by column
apply(P4b_data, 2, compute_proportion)
```

```
## [1] 0.6 0.6 0.6 0.6 0.6 0.6 0.6 0.6 0.6 0.6
```

It is not generating 10 independent samples but replicate one sample 10 times.

d

```
generate_flips <- function(prob) {
  rbinom(10, 1, prob)
}
probabilities <- matrix((31:40) / 100, ncol = 10)
P4d_data <- apply(probabilities, 2, generate_flips)
## calculate by row
apply(P4d_data, 1, compute_proportion)
```

```
## [1] 0.7 0.3 0.5 0.5 0.3 0.1 0.8 0.4 0.1 0.2
```

```
## calculate by column
apply(P4d_data, 2, compute_proportion)
```

```
## [1] 0.2 0.3 0.4 0.3 0.4 0.6 0.3 0.3 0.5 0.6
```

## Problem 6


```
dat6 <- readRDS("HW3_data.rds")
colnames(dat6) <- c("Observer", "x", "y")
```

1.

```
## coordinates values must be names as "x" and "y"
my_plot <- function(dat, title, xlab, ylab) {
  plot(dat$x, dat$y, main = title, xlab = xlab, ylab = ylab)
}
```


2.

```
## (a)
my_plot(dat6, "Single scatter plot", "dev1", "dev2")
```




```
## (b)
dat6_by_observer = split(dat6, f = dat6$Observer)
trash_can <- lapply(dat6_by_observer, my_plot, "Separate scatter plots", "dev1", "dev2")
```

### Separate scatter plots




## Separate scatter plots







## Separate scatter plots




### Separate scatter plots




## Separate scatter plots




## Separate scatter plots




## Separate scatter plots




## Separate scatter plots



## Separate scatter plots




## Separate scatter plots







## Separate scatter plots



## Separate scatter plots



## Separate scatter plots



## Problem 7

### Part a

```
#we are grabbing a SQL set from here
# http://www.farinspace.com/wp-content/uploads/us_cities_and_states.zip
#download the files, looks like it is a .zip
library(downloader)
download("http://www.farinspace.com/wp-content/uploads/us_cities_and_states.zip",dest="us_cities_states")
unzip("us_cities_states.zip", exdir=".")
#read in data, looks like sql dump, blah
library(data.table)
states <- fread(input = "./us_cities_and_states/states.sql",skip = 23,sep = "'", sep2 = ",", header = FALSE)
states <- states[-8,]
cities <- fread(input = "./us_cities_and_states/cities_extended.sql",sep = "'", sep2 = ",", header = FALSE)

### YOU do the CITIES
### I suggest the cities_extended.sql may have everything you need
### can you figure out how to limit this to the 50?
```

### Part b

```
## remove DC and PR
city_counts <- data.frame(table(cities$V4)[c(-8, -40)])
```

```
state_city_counts <- data.frame(city_counts, state = tolower(states$V2))
state_city_counts
```

```
##      Var1 Freq      state
## 1     AK  273     alaska
## 2     AL  838     alabama
## 3     AR  709     arkansas
## 4     AZ  532     arizona
## 5     CA 2651    california
## 6     CO  659     colorado
## 7     CT  438    connecticut
## 8     DE   98     delaware
## 9     FL 1487     florida
## 10    GA  972     georgia
## 11    HI  139     hawaii
## 12    IA 1060     iowa
## 13    ID  325     idaho
## 14    IL 1587    illinois
## 15    IN  989     indiana
## 16    KS  756     kansas
## 17    KY  961     kentucky
## 18    LA  725     louisiana
## 19    MA  703    massachusetts
## 20    MD  619     maryland
## 21    ME  489     maine
## 22    MI 1170     michigan
## 23    MN 1031     minnesota
## 24    MO 1170     missouri
## 25    MS  533     mississippi
## 26    MT  405     montana
## 27    NC 1090    north carolina
## 28    ND  407     north dakota
## 29    NE  620     nebraska
## 30    NH  284     new hampshire
## 31    NJ  733     new jersey
## 32    NM  426     new mexico
## 33    NV  253     nevada
## 34    NY 2207     new york
## 35    OH 1446     ohio
## 36    OK  774     oklahoma
## 37    OR  484     oregon
## 38    PA 2208    pennsylvania
## 39    RI   91     rhode island
## 40    SC  539    south carolina
## 41    SD  394     south dakota
## 42    TN  795     tennessee
## 43    TX 2650     texas
## 44    UT  344     utah
## 45    VA 1238     virginia
## 46    VT  309     vermont
## 47    WA  732     washington
## 48    WI  898     wisconsin
## 49    WV  859     west virginia
## 50    WY  195     wyoming
```

## Part c

```
letter_count <- data.frame(matrix(NA,nrow=50, ncol=26))
getCount <- function(letter, state_name){
  temp <- strsplit(tolower(state_name), "")
  count <- sum(temp[[1]] == letter)
  # how to count??
  return(count)
}
for(i in 1:50){ letter_count[i,] <- lapply(letters, getCount, states[i, 2]) }
```

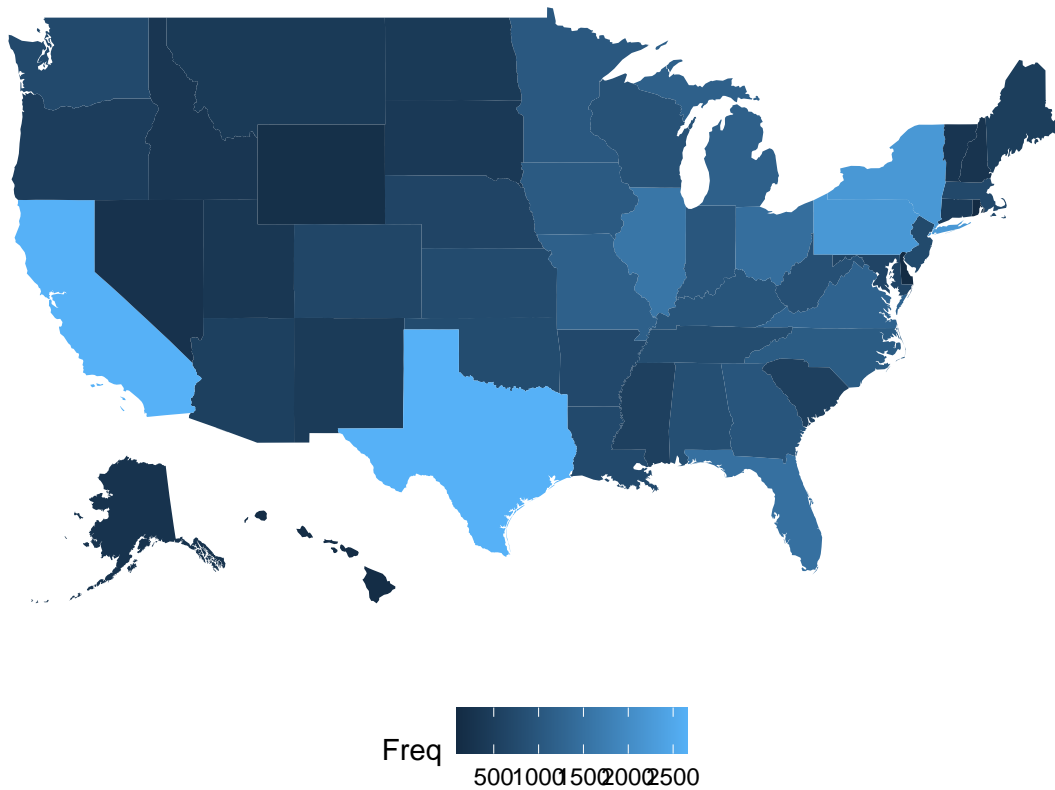
## Part d

```
#https://cran.r-project.org/web/packages/fiftystater/vignettes/fiftystater.html
library(ggplot2)
library(mapproj)
```

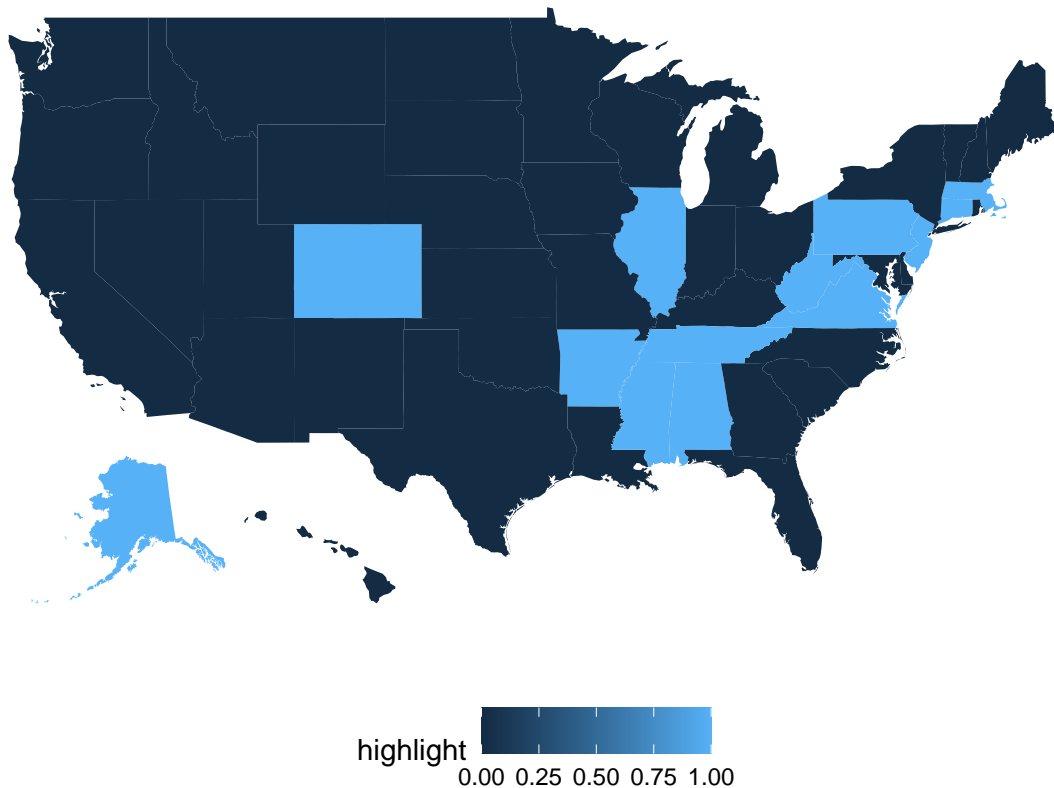
```
## Loading required package: maps
```

```
# crimes <- data.frame(state = tolower(rownames(USArrests)), USArrests)
# map_id creates the aesthetic mapping to the state name column in your data
load("fifty_states.rda")
p <- ggplot(state_city_counts, aes(map_id = state)) +
  # map points to the fifty_states shape data
  geom_map(aes(fill = Freq), map = fifty_states) +
  expand_limits(x = fifty_states$long, y = fifty_states$lat)+
  coord_map() +
  scale_x_continuous(breaks = NULL) +
  scale_y_continuous(breaks = NULL) +
  labs(x = "", y = "") +
  theme(legend.position = "bottom",
        panel.background = element_blank())
```

p



```
letter_highlight <- data.frame(state = tolower(states$V2),
                              highlight = apply(letter_count >= 3, 1,
                                                function(x){ifelse(sum(x) > 0, 1, 0)}} )
p2 <- ggplot(letter_highlight, aes(map_id = state)) +
  # map points to the fifty_states shape data
  geom_map(aes(fill = highlight), map = fifty_states) +
  expand_limits(x = fifty_states$long, y = fifty_states$lat) +
  coord_map() +
  scale_x_continuous(breaks = NULL) +
  scale_y_continuous(breaks = NULL) +
  labs(x = "", y = "") +
  theme(legend.position = "bottom",
        panel.background = element_blank())
p2
```



## Problem 8

```
library(tidyverse)
sensory_data_raw <- readRDS("sensory_data_raw.RDS")
sensory_data_tv <- sensory_data_raw %>%
  separate(col = "Operator", into = c("Item", as.character(1:5)), sep=" ", fill = "left") %>%
  fill("Item") %>%
  slice(2:n()) %>%
  gather(key = "Operator", value = "value", -Item)
```

### Part a

When he was “sd.boot[i]=coef(summary(lm(logapple08~logrm08, data = bootdata)))[2,2]”. He used logapple08 and logrm08, which are not the variable names of bootdata but existing vector.

### Part b

```
set.seed(3456)
Boot <- 100
coef_boot <- rep(0, Boot)
ind_box <- matrix(1:150, ncol = 5)
system.time({
  for(i in 1:Boot){
    # nonparametric bootstrap
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	time
apply	-0.329	-0.116	-0.035	-0.040	0.032	0.309	0.09
parApply	-0.298	-0.138	-0.052	-0.037	0.057	0.25	0.08

```
ind <- c(apply(ind_box, 2, function(x){sample(x, 30, replace = TRUE)}))
bootdata <- sensory_data_tv[ind,]
coef_boot[i]= coef(lm(as.numeric(value) ~ as.numeric(Operator), data = bootdata))[2]
}
})
```

```
##    user  system elapsed
##    0.08    0.00    0.07
```

## Part c

```
library(foreach)
library(doParallel)

#setup parallel backend to use many processors
cores <- detectCores()
cl <- makeCluster(cores-1) #not to overload your computer
registerDoParallel(cl)
set.seed(3456)
system.time({
  coef_boot_par <- foreach(i=1:100, .combine=cbind) %dopar% {
    ind <- c(apply(ind_box, 2, function(x){sample(x, 30, replace = TRUE)}))
    bootdata <- sensory_data_tv[ind,]
    coef(lm(as.numeric(value) ~ as.numeric(Operator), data = bootdata))[2]
  }
})
```

```
##    user  system elapsed
##    0.04    0.02    0.08
```

```
#stop cluster
stopCluster(cl)
```

Since bootstrapping samples are independent from each other, it can be completed on different computing node. That is why we can utilize parallel computing method.


I think in this particular question, they give similar result and their running time do not differ too much because it is a relatively small loop.

## Problem 9

```
x <- seq(-30, 2, 0.01)
y <- 3^x - sin(x) + cos(5*x)
plot(x, y, type = 'l', main = "f(x) = 3^x - sin(x) + cos(5x)")
```



$$f(x) = 3^x - \sin(x) + \cos(5x)$$



This function has countless roots.

## Part a

```
## target function
f <- function(x) {
  value <- 3^x - sin(x) + cos(5*x)
  return(value)
}

## derivative of target function
grad <- function(x) {
  value <- log(3) * 3^x - cos(x) - 5*sin(5*x)
  return(value)
}

find_the_root <- function(interval, tol) {
  ## tolerance is a single value while interval is a vector of dimension 2
  left <- interval[1]
  right <- interval[2]
  initial_point <- (left + right) / 2
  ## use while loop to solve the problem
  prev <- initial_point
  i <- 1
  while (abs(f(prev)) > tol) {
    ## calculate a new point
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	time
apply	-39.66	-31.15	-22.38	-22.51	-13.35	-4.97	129.64
parApply	-39.66	-31.15	-22.38	-22.51	-13.35	-4.97	20.64

```

curr <- prev - f(prev) / grad(prev)
## if the update goes beyond the given interval, assign limit to it
if (curr < left) {
  curr <- left
} else if (curr > right) {
  curr <- right
}
## update the previous point and draw it on the plot
prev <- curr
i <- i + 1
if(i > 1e06) break
}
## track progress
cat("Yes\n",file="log.txt", append = TRUE)
return(prev)
}

## try to find all the roots from -40 to -5
start_points <- seq(-40, -5, length.out = 1000)

start_intervals <- matrix(c(start_points - 1, start_points + 1), ncol = 2)

system.time({ roots <- apply(start_intervals, 1, find_the_root, 1e-06)})

```

## Part b

```

##setup parallel backend to use many processors
library(parallel)
cl <- makeCluster(cores - 1)
clusterExport(cl, "f")
clusterExport(cl, "grad")
system.time(roots_par <- parApply(cl, start_intervals, 1, find_the_root, 1e-06))
stopCluster(cl)

```

As we can see, the results both methods got are almost the same. And in this case, parallel computing is substantially faster than regular ones.